

1 Among conversation-level benchmarks, only this one re-derives its verdicts without the model.

	 The whole conversation judged as one unit	 Two people care recipient in scope	 What a reply leaves undone omissions scored	 Over time across sessions, memory	Replays verdict re-derives from saved answers, no model call
InvisibleBench this pilot · engine v2.0	●	●	●	●	●
K-Bench 2026 · 200 multi-turn vignettes	●	○	◐	○	○
HealthBench 2025 · 5,000 conversations, physician rubrics	●	◐	◐	○	○
CARE-Bench 2026 · simulated clients, psychological scales	●	○	○	○	○
RubRIX 2026 · 20,000 caregiver posts, one reply each	○	●	●	○	○
CounselBench 2025 · single-turn counseling, expert-rated	○	○	○	○	○
Single-turn safety and knowledge sets SafetyBench · HarmBench · EQ-Bench · ADRD-Bench	○	○	○	○	●
● documented ◐ partial ○ not documented. Read from each benchmark's own paper or README on 20 Sep 2026; readings and sources on the landing page. A design comparison, not an empirical one. This run exercises the first three properties; it is one session with no memory.					

THE FATIGUE CASE • CAREGIVER, TURN 1 OF 5

Five scripted caregiver turns; Mercury 2.5 replies live. Every check reads all five exchanges.

2026-09-19_14-04-44Z · InvisibleBench 5.1.0 · engine v2.0 · judge jev-43.0 · source ec1849e9 · Four scripted caregiver conversations, five turns each, Mercury 2.5 replying live · 40 judge requests · 2,391 saved answers · 196 derived judgments · \$0.016 · Madad, A. (2025). InvisibleBench: A Deployment Gate for Caregiving Relationship AI. arXiv:2511.20733 · HealthBench (arXiv:2505.08775) is related multi-turn work · No comparator, no ranking, no clinical outcome, no human labels, no held-out cases: the questions were written against saved conversations including these four.

Ask the same plan twice: 193 of 196 verdicts return. Change a rule: rerun for two cents. Every judgment, with its probabilities, is on the landing page.